

Can a photonic state-space model be trained on-chip? A pre-registered four-method bake-off on a realistic silicon-nitride ring substrate

Lucas Talandier — independent researcher, Paris

Abstract. Recurrent photonic processors promise low-latency, low-energy sequence processing, and the first optical recurrent network has recently been trained in situ — yet no continuous-time dissipative-resonator recurrence has ever had the parameters that define it — pole positions and inter-resonator couplings, the physics that *is* the memory — trained on the physical device by any gradient-based or gradient-estimating method (§1, §8). We ask whether such in-situ training is feasible for a structured photonic state-space model: an oscillatory (LinOSS-class) coupled-microring recurrence on ultra-low-loss silicon nitride. We (i) derive the mapping from the discrete oscillatory SSM to a physically realizable coupled-ring lattice and bound its realizable pole region; (ii) build one shared dissipative substrate model — finite Q, saturating gain, amplified-spontaneous-emission noise — with all parameters pre-registered before any training run; and (iii) run a four-method in-situ-training bake-off — SPSA, physics-aware training (PAT), recurrent in-situ adjoint, and Hamiltonian-echo learning (RHEL) — under a pre-registered fairness contract, scored on sample-efficiency-to-target-accuracy at matched device-pass cost. At realistic SiN noise, three of the four reach the pre-registered target on 8/8 seeds — both hardware-committed workhorses among them — with PAT needing $4.6\times$ fewer device passes than model-free SPSA (38,400 vs 176,000) but a digital-twin side-ledger four orders of magnitude more energy than SPSA’s entire training bill (13–44 J vs 2.6 mJ); RHEL is censored at budget, its echo defeated by the very dissipation the recurrence needs. A measured participation profile shows the single-drive lattice trains an effective dimension of ≈ 3 of $N=32$ rings, rising to all 32 with four input taps — a controllability constraint, and conversion cost, that any hardware implementation inherits. Trained damping is a first-order design knob ($\times 300$ in error across the feasible box), and training finds a heterogeneous damping profile that beats every uniform setting. An end-to-end envelope including electro-optic conversion and DAC/ADC overhead finds a conditional low-latency niche (GS/s line rates, $N \geq 32$, low-power heater class) — while an equally-calibrated offline-train-then-deploy baseline statistically ties in-situ training at calibration errors up to 30%-class; the tie breaks only under uncorrelated per-ring drift, where in-situ retraining holds a statistically significant, drift-growing edge that a global laser re-lock cannot absorb ($1.84\times$ time-integrated, below our pre-registered $2\times$ bar, so no formal advantage is declared): the demonstration is in-data; the advantage case is directional and drift-specific, and we say so. We release the pre-registration ledger, substrate model, and all training code.

1. Introduction

1.1 The gap

Photonic neural networks have learned to train themselves. In feedforward meshes and diffractive processors, on-chip training is now routine enough to have families: model-free perturbative methods [1, 2], hybrid physical-forward/digital-backward methods [3], and in-situ adjoint methods that read gradients from interference [4]. But the systems that most need on-device training — *recurrent* photonic processors, whose memory lives in the physics — have barely begun to receive it. Photonic reservoir computing deliberately avoids the problem: the recurrence is fixed, random, and only a readout is trained [5]. Where internal parameters of a photonic recurrence have been adjusted at all, it has been by calibration or regime-tuning rather than task-driven training [6], by reinforcement-style search over a readout while the loop stays fixed [7], or in systems whose recurrent state is carried digitally between optical passes [8]. A two-modality literature search with a pre-registered kill-criterion (§8; supplementary) found no demonstration, by any method, of the following:

a continuous-time dissipative-resonator recurrence — pole positions and inter-resonator couplings — trained in situ on the physical device by gradient-based or gradient-estimating methods on a computational task.

That sentence is this program’s target, with each qualifier load-bearing: *on a computational task* excludes the servo/calibration lineage; *physical parameters of the recurrence* excludes hybrid-digital state carriage; *weight-tied recurrence* excludes feedforward meshes folded in time. Refreshes of the search at assembly (2026-07-12) and a page-level verification round (2026-07-27; memos in supplementary) map the boundary against the strongest 2025–26 neighbors, which we dispatch by name because each is the “nearest miss” along one qualifier — and one of them moved the boundary. The **nearest neighbor** is the monolithic optical recurrent accelerator of Wu et al. [9]: its ORNN chip *is* trained in situ, by a model-free perturbative method (SPGD, two physical evaluations per update) on a classification task, with a weight-tied mesh applied across wavelength-encoded time steps — to our knowledge the first in-situ-trained optical recurrent network of any kind, and we cite it as such. What it does not do is train the parameters in the boxed sentence: its trained weights are interferometer-mesh voltages, its recurrent state is re-generated electronically at every step through a photodetector–modulator relay, and its resonators are calibrated once and held static — the continuous-time dissipative-resonator recurrence, whose *poles and couplings are themselves the memory*, remains untrained. The remaining near-misses each fall along one qualifier: on-chip all-photonic backpropagation is now demonstrated — for a *feedforward* network [10]; microring weight banks have been trained in situ through on-chip optical backprop — as *feedforward* layers [11]; a time-synthetic fiber-loop network trains in situ — with per-step distinct programmed parameters, i.e. unrolled feedforward rather than a weight-tied recurrence [12]; a modulation-and-weighting microring array trains weights sitting *inside* an analog recurrent loop on-device — by particle-swarm search, explicitly without gradients [13]; an optoelectronic delay reservoir has had recurrence-defining parameters optimized in situ — by Bayesian search, through a digital feedback loop [14]; and a silicon photonic reservoir equalizer is “trained in hardware” — its readout only, by an evolution strategy, the reservoir couplings fixed by fabrication [15]. No coupled-resonator lattice has had its couplings learned on-device by any method. (The claim is time-indexed: the search refresh re-runs once more immediately before submission, with the registered page-level reads listed in supplementary.)

1.2 Why a state-space model, and why silicon nitride

Two developments make this the right moment to close the gap. On the algorithmic side, deep state-space models showed that *linear* recurrences with well-placed poles — not gated nonlinear dynamics — carry most of long-sequence performance [16, 17], and the oscillatory LinOSS line extended this to second-order units that are exactly coupled damped oscillators [18, 19]. A lattice of coupled microrings *is* such a system: poles are ring detunings and losses, couplings are physical couplers, and — decisively for hardware — the LinOSS stability analysis tolerates nonnegative damping, so a lossy-but-high-Q dissipative lattice is in-regime rather than an approximation to a conservative ideal (§2). Damping becomes a design knob, not an embarrassment.

On the platform side, the figure of merit for a dissipative ring recurrence is memory per pass — round-trip loss, hence intrinsic Q. Ultra-low-loss silicon nitride is class-leading and foundry-accessible (Q_i near 10^7 on multi-project-wafer runs [20]), thermally stable enough for the slow thermo-optic actuation that gradient-estimating training wants, and needs so little gain that the injected amplifier noise stays low (§3). The same choice has a cost we do not hide: SiN’s weak $\chi^{(3)}$ makes the one training route that needs an optical nonlinearity — Hamiltonian-echo learning, whose echo is phase conjugation — *harder*, not easier, and §4–§5 treat that honestly rather than idealizing it away.

1.3 What this paper does

This is a theory-and-simulation (Stage-0) paper, built to be falsified before it is fabricated. Its results are gated by a pre-registration ledger — margins, budgets, fairness conventions, and statistical rules frozen (and PI-signed) before the runs that consume them, with the commit trail shipped as supplementary material (§3.7). Concretely:

1. **A mapping with a realizable region** (§2): the oscillatory-SSM $\leftrightarrow$ coupled-SiN-ring correspondence, with the pole region bounded by measured platform numbers, an actuation map for $\{\delta, \kappa_{\text{ext}}, \mu\}$, and the backscatter doublet carried as the model rather than a correction.
2. **A shared dissipative substrate** (§3): saturating Er gain at a registered operating point, Langevin ASE at NF 7 dB, three cells from foundry-floor to aspirational, a measured feasible box bounded by the lasing crossing, and a measured minimal multi-point input map after single-drive controllability was found to collapse (effective dimension ≈ 3 of 32 rings).
3. **A four-method bake-off under one fairness contract** (§4–§5): SPSA, PAT (with registered twin-mismatch), a recurrent in-situ adjoint (charged as if realizable; optimistic bound, labelled), and RHEL over a concrete χ^3 -FWM echo sub-model — scored by device passes to a pre-registered target relative to the exact-gradient ceiling, eight seeds, censoring-aware statistics.
4. **The result** (§5): both hardware-committed methods train the recurrence to within margin of the ceiling on all eight seeds — the sentence in §1.1, demonstrated in simulation at a foundry-class noise cell. The ranking (PAT < adjoint < SPSA on device passes; RHEL censored, and worse than readout-only under an honest echo) settles the hardware roadmap on PAT/SPSA without promotion of the exotic routes. And the comparison the fair design was built to expose lands as a sharpened null with one directional exception: calibrate-then-deploy ties in-situ training at calibration errors all the way to 30%-class — calibration accuracy is not the axis on which in-situ training pays — while uncorrelated per-ring drift, which a global laser re-lock cannot absorb, gives in-situ retraining a statistically significant edge that grows with accumulated drift yet stays below our pre-registered $2\times$ advantage bar, so no formal advantage is declared (§5.5, §7).
5. **Limits, stated as limits** (§8): robustness is measured against modelled imperfections only; the anchor risks, verification debts, and the single-session review period are disclosed with the same specificity as the results.

Our position on novelty is deliberately narrow. PAT and SPSA are chip-proven; we do not re-validate them. What has never existed is a *recurrent, dissipative* photonic system trained through its own physics — and a demonstration that the recurrence-defining parameters of a realistic SiN lattice can be so trained, under pre-registered thresholds and honest costing, is the contribution. Whether it *pays* is a separate question (§7), and this paper reports the current answer to that question as it falls, not as we might wish it.

2. From oscillatory state-space models to coupled SiN microrings

2.1 The model class

A structured state-space model (SSM) processes a sequence $\mathbf{u}_1, \mathbf{u}_2, \dots$ through a linear recurrence $\mathbf{x}_{t+1} = \bar{\mathbf{A}}\mathbf{x}_t + \bar{\mathbf{B}}\mathbf{u}_t$, $\mathbf{y}_t = \mathbf{Re}(\bar{\mathbf{C}}\mathbf{x}_t) + \mathbf{D}\mathbf{u}_t$, whose expressive power is set by where the eigenvalues (poles) of $\bar{\mathbf{A}}$ can be placed and how precisely. The diagonal variants that dominate current practice — S4D and DSS [17, 21] — reduce $\bar{\mathbf{A}}$ to a set of independent complex poles, $\mathbf{H}(\mathbf{s}) = \sum_j \mathbf{c}_j \mathbf{b}_j / (\mathbf{s} - \lambda_j) + \mathbf{D}$; the oscillatory LinOSS family [18] instead parameterizes forced, damped harmonic oscillators, and its damped extension D-

LinOSS [19] makes the damping of each mode a trainable parameter. Crucially for what follows, the LinOSS construction is valid for any nonnegative-diagonal (dissipative) state matrix: stability does not require conservative dynamics, only that every mode decay. A physical substrate that is *lossy but slowly so* is therefore in-regime by construction — the loss is not an error term to be fought but the damping parameter of the model class itself.

2.2 One ring is one trainable complex pole

The mode amplitude $\mathbf{a}_j$ of a silicon-nitride microring obeys temporal coupled-mode theory (Haus energy-amplitude convention, $|\mathbf{a}|^2 = \text{stored energy}$) [22]:

$$\dot{\mathbf{a}}_j = (i\delta_j - \kappa_{\text{tot},j}) \mathbf{a}_j + i \sum_{k \neq j} \mu_{jk} \mathbf{a}_k + \sqrt{2\kappa_{\text{ext},j}} u(t),$$

with $\kappa_{\text{tot},j} = \kappa_{i,j} + \kappa_{\text{ext},j}$ the amplitude decay rate (intrinsic + external), δ_j the detuning of the drive from the ring resonance, and μ_{jk} the inter-ring coupling. A single uncoupled ring is therefore exactly one **complex** Laplace pole,

$$\mathbf{s}_j = -\kappa_{\text{tot},j} + i\delta_j,$$

— complex, not real, because the optical envelope carries a carrier. A bank of N rings natively realizes a **diagonal complex-pole SSM of the S4D/DSS class**: one ring per pole, with the injection and readout couplings playing the role of the residues $\mathbf{c}_j \mathbf{b}_j$. The oscillatory LinOSS unit is recovered *exactly* as the uncoupled, real-input/output special case: a real damped second-order oscillator is a conjugate pole pair, and the identification against the LinOSS eigenvalue $\mathbf{a}_i = -e^{\alpha_i} + i\beta_i$ is

$$e^{\alpha_i} = \kappa_{\text{tot}}, \quad \beta_i = \delta.$$

Storing $\alpha_i = \log \kappa_{\text{tot}}$ enforces the dissipative sign and $\kappa_{\text{tot}} > 0$ by construction — the physical substrate implements the model’s own stability parameterization. We deliberately claim the broader class and treat LinOSS as its special case, with one consequence made explicit now: published LinOSS/D-LinOSS benchmark results validate only the $\mu = 0$ diagonal reduction, and transfer to the coupled realization is verified in-house against a BPTT-on-substrate ceiling (§5.1) rather than assumed.

Discretization is exact, not approximate: for piecewise-constant input the zero-order-hold (van-Loan matrix-exponential) step gives discrete poles $\mathbf{z} = e^{\mathbf{s} \, dt}$ to machine precision, with dt the round-trip time. The per-step memory retention is $|\mathbf{z}| = e^{-\kappa_{\text{tot}} dt}$; long memory is literally $|\mathbf{z}| \rightarrow 1$.

Validation. The mapping is validated at three independent levels (test-anchored; all thresholds pre-specified): (i) *by construction* — uncoupled system poles match the CMT reference to $< 10^{-3}$ and discrete $|\mathbf{z}|$ to $< 10^{-9}$; (ii) *in the CW limit* — the dynamical model’s steady state recovers independently derived static transfer functions with error scaling as $O(1/\mathcal{F})$ (finesse), measured 4.5×10^{-4} at $\mathcal{F} = 1048$ falling to 4.4×10^{-5} at $\mathcal{F} = 10473$; the registry rings sit at $\mathcal{F} \approx 10^3 - 1.5 \times 10^4$, where the CMT pole is an excellent model; (iii) *in the time domain* — against an independent RK45 integration, ringdown matches the closed form to $< 10^{-9}$ and both pole coordinates (κ and δ) re-fitted from the trajectory recover the set values to $< 10^{-5}$.

2.3 Coupling rings: the structured off-diagonal generalization

Energy-conserving inter-ring coupling enters as $i\Omega$ with Ω real-symmetric — an anti-Hermitian addition to the state matrix. It hybridizes the poles (photonic-molecule supermodes) while conserving total damping: the eigenvalues move along the imaginary axis while $\sum_j \text{Re } \lambda_j = -\sum_j \kappa_{\text{tot},j}$ is invariant (verified: two-ring beat

frequency equals the eigenvalue splitting 2μ to $< 2\%$, with a no-beat control at $\mu = 0$). The trainable μ_{jk} are thus a *mild, structured* generalization beyond diagonal- $\mathbf{A}$ SSMS — off-diagonal terms that reshape frequencies but cannot buy stability or memory. They are also load-bearing for the claim this paper builds toward: inter-ring couplings are recurrence-defining parameters (§2.5).

2.4 The realizable pole region

Where can the poles actually go? Four physical boundaries define the envelope (Fig. F1).

Memory is loss-limited; stability is free. Passive rings satisfy $\kappa_{\text{tot}} \geq \kappa_i = \omega_0/(2Q_i)$, so the maximum amplitude (state) memory time is $1/\kappa_i = 2Q_i/\omega_0$ — twice the photon-energy lifetime, a convention this paper uses consistently. Across the registered platform range this gives **3.29 ns (329 round trips)** at the foundry-conservative corner ($Q_i = 2 \times 10^6$) to **49.4 ns (4937 round trips)** at the class-leading corner ($Q_i = 3 \times 10^7$): a $15\times$ memory span that is the platform’s dominant figure of merit and the reason this program is on SiN (§1). Net optical gain moves $\kappa_{\text{tot}} \rightarrow 0$ and hence $|z| \rightarrow 1$, extending memory past the passive floor at the cost of saturation and amplified-spontaneous-emission noise — deferred to the substrate model (§3).

Frequency is FSR-bounded. The detuning axis aliases at one free spectral range: $|\beta_j dt| \leq \pi$ with $dt = 1/\text{FSR}$. The reachable imaginary axis is one FSR wide.

Readout trades against memory. The external coupling sets both memory ($1/\kappa_{\text{tot}}$, maximized by deep undercoupling) and readability (residue $\propto 2\kappa_{\text{ext}}$; on-resonance drop efficiency $(2\kappa_{\text{ext}}/\kappa_{\text{tot}})^2$, i.e. detector SNR) in opposite directions; their product is bounded by the passive memory length (tested). At the foundry corner the trade runs from 274 round trips at drop efficiency 0.028 (deep undercoupling, $\kappa_{\text{ext}} = 0.1\kappa_i$) to 16 round trips at 0.91 (deep overcoupling, $10\kappa_i$); the same Pareto shape holds at the class-leading corner with the memory axis scaled $\sim 15\times$. The operating point is not chosen here: it is pre-registered (§3, PR-4), and κ_{ext} is itself a pole-real-part actuator, so a trainable- κ_{ext} policy folds the readout trade into the recurrence training.

Backscatter bounds the one-ring-one-pole abstraction — by roughness, not cleanly by Q . Sidewall roughness couples the counter-propagating modes at a coherent rate γ that is a fabrication constant, essentially independent of κ_{tot} ; the resonance resolves into a standing-wave doublet when $2\gamma \gtrsim \kappa_{\text{tot}}$ (conservative HWHM criterion; the FWHM criterion shifts every crossover $\times 2$, conclusion unchanged) [23]. Because κ_i falls as Q_i rises while γ does not, splitting *grows* with Q : from measured SiN rates, a clean damascene-class process ($\gamma/2\pi \approx 12$ MHz [24]) keeps the single-pole picture at the foundry corner ($2\gamma/\kappa_i \approx 0.5$) and breaks it at $Q \approx 4 \times 10^6$ (HWHM; 8×10^6 FWHM), while a rough subtractive process splits 21–75 % of resonances *already* at $Q_i \lesssim 2.7 \times 10^6$, with average doublet separations of 180–320 MHz depending on etch mask [25] — i.e. modal-coupling rates $\gamma/2\pi \approx 90\text{--}160$ MHz under the standard 2γ -separation convention [23], a mapping we state here because the source tabulates separations, not rates (page-verified 2026-07-12). *The assembled crossover curve brackets published SiN data points — no single published SiN crossover exists, and no γ is published for the specific target processes; primary-source verification is registered before submission.* Two consequences propagate forward: the substrate model carries a roughness-gated CW/CCW doublet knob, default ON except at the clean-damascene corner (§3); and a platform tension is on record — the best-memory (highest- Q) platforms are the most splitting-prone, while the splitting-safe low- Q foundry corner is memory-poor. Overcoupling suppresses the visible doublet, so the κ_{ext} policy and the splitting risk are coupled: the max-memory (deep-undercoupled) regime is the most exposed.

2.5 What “training the recurrence in situ” means physically

The recurrence is defined by the state matrix $\mathbf{M} = \text{diag}(-\kappa_{\text{tot},j} + i\delta_j) + i\Omega$. Every entry has a concrete, independently addressable actuator:

Pole quantity	Physical knob	Actuator	Cost
Im (frequency) δ_j	ring resonance offset	thermo-optic heater (trench-isolated), one per ring	low: 1 heater + DAC channel/ring
Re (damping) $\kappa_{\text{tot},j}$	bus–ring coupling $\kappa_{\text{ext},j}$	tunable coupler (MZI-assisted gap)	medium: 1–2 heaters/ring
same, past the passive floor	per-ring net gain g_j	Er:SiN / III–V pump current	high (gain stage + ASE cost); not required for the claim
off-diagonal μ_{jk}	ring–ring coupling	tunable photonic-molecule coupler (or bus-mediated mesh)	medium–high; topology fixed at fab, strengths trainable
residues B, C	injection/readout mesh	thermo-optic MZI mesh	not recurrence-defining

The last row is the load-bearing exclusion. Training only the injection/readout mesh — the residues — while the poles stay fixed is precisely reservoir computing with a trained linear head, and prior photonic work in that regime (including reinforcement-learning-tuned readouts [7]) does not train the recurrence. The partition this program registers (§4, PR-2/PR-6) is that a method earns the in-situ-training claim only if it updates the pole-defining set $\{\delta_j, \kappa_{\text{tot},j}, \mu_{jk}\}$ on the device, and every method in the bake-off trains the same partition. Notably, the minimal set is *gain-free* and all-thermo-optic — heaters and tunable couplers only — which is what makes it drift-stable and realistic for perturbative (SPSA-class) training on SiN; no active gain stage is required for the headline claim.

The mapping, then, is not an analogy but an identification: a coupled SiN ring bank *is* a diagonal-complex-pole SSM with structured coupling, its loss *is* the model’s damping parameter, its realizable pole region is bounded and known, and every recurrence-defining parameter has a physical actuator. Whether those parameters can be *trained through the physics* — with realistic gain saturation, ASE noise, and measurement cost — is the question the rest of this paper is built to answer.

3. A pre-registered dissipative substrate

3.1 Why one shared substrate

The bake-off’s conclusions are comparisons, and comparisons inherit the honesty of their common ground. All four training routes therefore run on a *single* simulated substrate — one dynamical model, one noise model, one parameter partition, one initialization and data convention (PR-6) — frozen and signed before any estimator existed. Nothing in the substrate was chosen with knowledge of which method it would favor; the freeze order is auditable in the ledger.

The model is the 2N-mode coupled-mode-theory (CMT) lattice of §2: N rings, each carrying a clockwise/counter-clockwise doublet (backscatter coupling is always on, §3.4), with the linear readout followed by direct detection $y(n) = |\sum_j c_j a_j(n)|^2$ — the substrate’s only nonlinearity apart from gain saturation. The trainable (in-situ) partition is fixed by the architecture freeze (PR-2): per-ring detunings δ_j , per-ring tunable bus couplings $\kappa_{\text{ext},j}$, and nearest-neighbor inter-ring couplings $\mu_{j,j+1}$ — the parameters that *define* the recurrence, as opposed to reservoir-style approaches that train only the readout. The digital head (an 8-lag linear FIR on y) and the fixed affine intensity encoder bracket the photonic core; both are identical across methods, and the head’s training cadence is part of the fairness contract.

3.2 Loss, gain, and the operating point

The per-ring net decay is $\kappa_{\text{net},j} = \kappa_i - g_j + 2\kappa_{\text{ext},j}$: intrinsic loss, minus Er gain, plus the loading of the through/drop ports. Gain follows the **M1 static-saturated** class: a per-episode operating point $g(\bar{P}) = g_0/(1 + \bar{P}/P_{\text{sat}})$, differentiable in the drive statistics and in κ_{ext} (the coupling changes the intracavity power the medium sees), fixed within a rollout — valid under the registered conditions (stationary within-episode drive statistics, pinned train/test power statistics, quasi-static margin $E_{\text{sym}}/E_{\text{sat}} \sim 10^{-6} - 10^{-5}$, all [EV]-anchored; bursty inputs void M1). Every gain-bearing cell runs at the registered ceiling $g_{\text{rt}} = 0.9 \times$ the intrinsic round-trip loss — gain never compensates port loading — so $\kappa_{\text{net}} = 0.1\kappa_i + 2\kappa_{\text{ext}}$ at the operating point, with the passive floor ($g = 0$) as a labelled sensitivity row. A one-off rate-equation integrator (the salvaged M2) validated the quasi-static reduction at the gating cells.

Amplified spontaneous emission enters as a continuous Langevin term (convention A2), $\langle FF^* \rangle = 2\kappa_g n_{\text{sp}} \delta(t-t')$ in photon units, discretized exactly alongside the ZOH/van-Loan propagator, with **fresh draws on every physical pass** — noise is never shared between passes, methods, or the forward/reverse directions (this single convention carries much of the fairness burden, and one of RHEL’s irreversibility invariants). The headline noise figure is **NF = 7.0 dB** ($n_{\text{sp}} = 2.5$). The program’s original debt register carried “no measured NF for the flagship Er:Si₃N₄ device” as verification debt #3; the S0.3-0 substrate reconnaissance (2026-06-10) found that premise false — the full text reports “a noise figure of ca. 7 dB ... at net gain of >20 dB, limited by coupling losses” [26] — and NF-A = 7.0 dB was frozen *to match that measurement* (re-verified at page level during assembly, 2026-07-12). The residue of the debt is narrower than its original form: what exists is a single coupling-loss-limited *system* NF, not an isolated intrinsic amplifier NF, so $\text{NF} \in \{3, 5\}$ remain registered sensitivity values.

3.3 The three cells

cell	platform class	Q_i	N	role
C-1	foundry floor (P-FND)	2×10^6	8	Gate-ii gating cell
C-2	demonstrated MPW class (P-AN800)	6.8×10^6	32	bake-off headline
C-3	aspirational (P-UHQ)	3×10^7	128	labelled sweep axis; never gates

C-2 is registered as a *conservative bound* on a demonstrated MPW result (3.3 dB/m, mean $Q_i \approx 10.8\text{M}$ [20]); our pair (5.1 dB/m, $Q_i = 6.8 \times 10^6$) is strictly worse than the broadest-linewidth device in that paper’s full 249-resonance distribution, whose figures are independently assessed in an invited peer commentary [27]. The residual anchor risk is geometry transfer (the demonstrated loss is achieved *by* a wide multimode Euler-bend racetrack; our registry ring is single-mode) — flagged by the commentary itself and priced by a registered $\times 2$ -loss derate row. At the operating point, C-2 holds ~ 32 samples of memory at 2 GS/s and packs the 32-ring spectrum in-band; C-1 holds ~ 9.4 (covering the task’s 7-tap span at the operating point; marginal only at its passive floor).

3.4 Backscatter and the doublet

Surface roughness couples the CW/CCW modes at rate γ (11.8 MHz at C-2, damascene-clean process class; 90 MHz at C-1), splitting each ring resonance into a doublet — at high Q_i the splitting is *resolved* ($\gamma/\kappa_{\text{net}} \approx 16.6$ at C-2/ θ_0), so the substrate is the full **2N**-mode doublet lattice **always**, not an N-mode idealization with a correction term. This choice was consequential: the S0.4-0 calibration measured the drive build-up at the lasing floor to be $\times 0.42$ of the hold value — the doublet plus chain hybridization *quench* the single-pole build-up that two independent prior estimates ($\times 37$ and $\times 91$) had presumed, mooting both. Modelled imperfections cannot be un-modelled by a reviewer; unmodelled ones are the subject of §8.

3.5 The feasible box, the lasing boundary, and the drive budget

The trainable couplings live in a registered box $\mathbf{r}_j = \kappa_{\text{ext},j}/\kappa_i \in [0.1, 3]$ (K4), verified to map into the realizable pole region at every cell. With saturating gain the box's lower edge is unsafe: below $\mathbf{r}^* \approx 0.134$ the ring crosses threshold and the linear rollout diverges. The operative bound is the S0.4-0-measured $\mathbf{r}_{\min} = 0.1606$ (crossing + registered margins), enforced as the training-time clamp; the θ_0 hold value is $\mathbf{r}_0 = 0.3$. Damping, in the D-LinOSS sense, is deliberately *not* a frozen cell: it is the trainable per-ring net loss the estimators explore through κ_{ext} over this box at fixed $\mathbf{g}_{\text{rt}}$ (the R-ii disposition), characterized separately (§6).

Drive is normalized by an intracavity-energy budget (O2): the registered $\bar{P}_0 = 1$ mW bus drive corresponds to $E_0 = 1.069 \times 10^8$ intracavity photons at θ_0 , and the digital encoder is *not permitted to buy SNR* by rescaling drive power against the registered noise cell — the budget is held fixed across methods and across the in-situ/offline comparison. The S0.4-0 calibration confirmed the budget is invariant (**1.000000**) under the four-tap input map below.

3.6 The measured input map, and what it says about controllability

A single bus port cannot train a 32-ring chain: with drive on ring 1 only, the per-ring gradient magnitude collapses with hop distance — an effective participating dimension of ≈ 3 of 32 rings. The registered remedy is a measured, minimal multi-point input map: S0.4-0 searched tap sets under a pre-registered every-ring controllability gate (each ring's gradient $\geq 10^{-3}$ of the maximum, min over five drive seeds) and resolved $\mathbf{B} = \text{taps } \{3, 12, 21, 30\}$ at $\mathbf{K} = 4$ (32/32 rings pass, worst 1.42×10^{-3} ; no $\mathbf{K} \leq 3$ set passes). The four E/O drive channels this costs are charged to the systems envelope (§7) — multi-point drive raises exactly the conversion overhead the envelope exists to price. Two protocol rulings made during the measurement (gate referenced to the maximum ring; min-over-seeds robustness) both strengthen the gate and are recorded for review. One anchor risk stays open (vii): the calibration's floor numbers are computed on-resonance; a detuned ring at $\mathbf{r}_{\min}$ can reach $\kappa_{\text{net}} = -0.48\kappa_i$ under worst-case de-saturation, and a δ -aware floor would sit at **0.2547** — the conformance check that would settle it was measured to be not-cheap and is carried as a labelled risk, not silently absorbed.

3.7 The ledger as method

Every consequential choice above — the gain class and its validity conditions, the NF axis, the cells, the box, the clamp, the input map, the drive budget — was written into a pre-registration ledger, in most cases signed by the PI, *before* the run that consumed it, with measured deferrals returned to the ledger as dated addenda (the E_0 value, $\mathbf{r}_{\min}$, the resolved input map, the budget $\mathbf{B}$, the BPTT ceiling). Supersessions retain the superseded text. The ledger, its review trail (independent adversarial review through 2026-07-06; single-session self-review thereafter, disclosed), and the frozen-before-run commit hashes ship as supplementary material. We treat this as part of the method: a bake-off whose thresholds, budgets, and fairness conventions are set after seeing results would not support the claims of §5.

4. Four routes to on-chip gradients

4.1 What counts as a physical gradient

The fairness contract (PR-6) fixes one invariant above all: **a training method may obtain gradient information only from simulated device passes on the shared substrate, with fresh noise on every pass.** Autodifferentiation through the substrate is reserved for the BPTT reference — the ceiling, never a contestant. All methods share, per seed: the same initialization (detunings spread over $[-\kappa_i, \kappa_i]$, $\mathbf{r}_0 = 0.3$, connected chain

$\mu_c = 0.3\kappa_i$), the same data stream (a function of seed and iteration only), the same head cadence, the same clamp to the operative box, and the same pre-registered hyperparameters. Cost is counted in **physical device passes, any direction** (PR-7): one pass = one sequence through the substrate. Digital compute (a twin’s forward/backward, the head update) lives on a side-ledger that is always co-reported and never folded into the rank — a method that saves device passes by spending digital FLOPs has a real but *different* advantage than a model-free one, and merging the ledgers would hide exactly the distinction the comparison exists to draw.

4.2 SPSA — model-free, two passes

Simultaneous-perturbation stochastic approximation [1] perturbs the entire in-situ partition by $\pm c\Delta$ (a random sign vector), measures the scalar loss twice, and forms a descent direction from the difference: **2 device passes per update**, no model, no twin, no added hardware beyond the plant’s own actuators and its single readout (the simplest row of the hardware ledger, F8). Perturbations at the clamp boundary are one-sided. SPSA is chip-demonstrated [2] and inherits the crosstalk-robustness observed in our prior thermo-optic work [28]; its known weakness — gradient variance growing with parameter count — is precisely what the sample-efficiency metric prices.

4.3 PAT — physical forward, twin backward

Physics-aware training [3] evaluates the loss on the *measured* physical output and takes the parameter gradient through a differentiable digital twin at the commanded parameters: **1 device pass per update**, plus a twin forward/backward on the digital ledger. PAT’s value proposition is exactly this exchange, and its honesty hinges on the twin being *imperfect in a registered way*. The twin-mismatch protocol (PR-5) freezes three families: **M-par** — 5%-class parametric calibration errors on the constants and actuation maps ($\kappa_i+5\%$, $\gamma+5\%$, actuation $\times 1.05/\times 0.95$, detuning offset $0.05\kappa_i$, a loose gain pair $\times 1.10/\times 0.75$ reflecting debt #3); **M-struct** — a structural omission (the twin drops the gain’s dependence on the trained coupling, the one channel a fixed-gain model cannot see); and **M-noise** (always on) — the twin is noiseless. The headline PAT is the *composed* mismatch; the decomposition is reported separately (§5.6). A perfect-twin gate verifies the implementation: with mismatch off, PAT’s gradient equals BPTT’s to machine precision.

4.4 Recurrent in-situ adjoint — a physical reverse pass, by hypothesis

The adjoint route extends feedforward in-situ backpropagation [4] to the time-domain cavity setting: the error field is physically propagated *backward* through the same dissipative substrate, and the gradient is read from forward/adjoint interference. No recurrent photonic demonstration of this pass exists (verification debt #4); the simulation charges it **as if realizable** — 1 forward + 1 adjoint = **2 device passes per update**, zero digital — and quarantines the realizability question in the hardware ledger (circulators, phase-coherent injection, separability of the counter-propagating field from the backscatter doublet). The simulated adjoint pass is honest about two physical limits: it carries **fresh noise** (the reverse pass is its own noisy traversal), and the saturating gain is **frozen at its operating-point value** in the backward linearization — a counter-propagating field experiences the medium’s saturation state but cannot realize the $\partial g/\partial \kappa_{\text{ext}}$ self-consistency channel. Both simplifications flatter the adjoint (a third, an additive noise term on the adjoint field itself, awaits an error-launch power convention that is unresolved hardware design), so the bake-off’s adjoint arm is an **optimistic bound** and is labelled as such wherever it appears. Floor checks: on the fixed-gain plane the adjoint gradient equals BPTT to machine precision; in saturating mode the frozen-gain approximation costs 0.6% of gradient direction at C-1 but **7.5% at C-2** — the one method-relevant quantity we measured to *grow* with cell size.

4.5 RHEL — Hamiltonian echoes on a substrate that forgets

Recurrent Hamiltonian echo learning [29, 30] trains by time-reversal: evolve forward; apply a single conjugation to the state snapshot (for optical fields, phase conjugation); evolve again through the *same* physics with the input replayed time-reversed and a small error nudge injected continuously; the gradient is the symmetric finite difference of $\nabla_{\theta} H$ between a $+\epsilon$ and a $-\epsilon$ echo. Three consequences of taking the published algorithm seriously on a dissipative substrate, each registered before the runs: (i) the count is **3 passes in the Hamiltonian limit but 4 operationally** (2 forward + 2 echo) — the echo does not return the state re-usable, and state cloning is unphysical (PR-7.1); (ii) the conjugation fires **twice per update**, paying its penalty chain independently each time; (iii) the update rule reads only the *coherent* generator, so the dissipative channel of κ_{ext} is structurally invisible to it.

The echo’s physical primitive is modelled concretely (PR-11), not as an idealized operator: a $\chi^{(3)}$ four-wave-mixing conjugation stage with the full penalty chain — extraction through the ring ports ($\eta_{\text{ex}} = 2\kappa_{\text{ext}}/\kappa_{\text{net}} = 0.86$ at θ_0), single-pass spiral conversion ($(\gamma_{\text{nl}} P_p L)^2 \approx -16.7$ dB at 0.3 W pump, 0.5 m; $\gamma_{\text{nl}} \approx 0.97 \text{ W}^{-1} \text{ m}^{-1}$ for tight-confinement SiN [31]), and timing decay — totalling **-22.4 dB per conjugation** (Fig. S3), plus the phase-insensitive parametric quantum floor. A page-level check at assembly found published *ultra-low-loss-geometry* demonstrations at $\gamma \approx 0.29\text{--}0.51 \text{ W}^{-1} \text{ m}^{-1}$ (with CW power handling demonstrated to 7 W) [31]; our 0.97 assumes a tighter-confinement spiral than those demos, so the frozen chain is, if anything, *optimistic* for RHEL — at the measured ULL values the conjugation penalty deepens by a further ~6 dB. The direction only strengthens §5.4’s conclusion, which the idealized-conjugator control shows does not hinge on the chain at all. The quantum floor itself was measured negligible at the 10^5 -photon state scale (as registered-to-measure), with pump-transfer excess included. Because the ring fields overlap spectrally, conjugating N rings needs N pumped arms: **≈ 9.6 W of on-chip pump at the headline cell**, charged to the envelope. Off-chip conjugation was priced and excluded (the state decays in transit: amplitude survival 0.12–0.54 at 10 ns for C-1/C-2; there is no storage primitive to wait out a conjugator). Three irreversibility invariants bind the implementation and are test-enforced: independent forward/echo noise streams (no common-RNG reversal); no loss-sign flip (the echo traverses the same dissipative lattice); gain injects fresh ASE in the echo too — and the echo of a noisy forward must *not* recover the noiseless state. The floor check completes the picture: as dissipation is removed, the implemented RHEL gradient converges to the exact reference (cosine $\rightarrow 1.0000$), so whatever §5.4 finds is the physics, not the code.

4.6 The guardrail

The four routes are *parallel in simulation, singular in hardware*: the Stage-1 chip is committed to the two chip-demonstrated workhorses (PAT, SPSA) regardless of the bake-off’s ranking, and an exact-gradient route can earn a *later* hardware slot only by clearly beating both on a pre-registered outcome criterion (PR-9; “exactness” is struck from the promotion menu). §5 reports how this resolved.

5. The bake-off: which routes train the recurrence, and at what cost

5.1 Pre-registration and the reference ceiling

Every threshold in this section was fixed in the ledger *before* the run it judges, in the order the runs consumed them: the fairness contract, cost metric, and twin-mismatch families (PR-6/7/5) at the start of S0.4; the target rule, statistical plan, and gate semantics (PR-3/8/9) at S0.4-close; and the two measured numbers a rule cannot supply — the budget and the ceiling — as dated addenda committed before any contestant ran. We state this not as

ceremony but because the central claim is a *threshold-crossing* claim, and a threshold chosen after seeing the curve is worthless.

The reference is backpropagation-through-time on the substrate itself (BPTT), which is not a physical training method — it reads gradients the device cannot expose — but bounds what the task admits at this cell. On the headline cell (C-2: 32 rings, $Q_i = 6.8 \times 10^6$; 4-PAM channel equalization at 28 dB, §3), BPTT drives the symbol-error rate to a median of 5.2×10^{-4} across eight seeds (seven of eight at 5×10^{-4} — two errors in the 3840-symbol evaluation set, the quantization floor). **The frozen substrate has ample capacity for the task; the open question is purely which physical training routes reach it, and at what cost.** The pre-registered target follows mechanically: a method *reaches target* if its held-out SER falls to $\text{SER}_{\text{target}} = 1.25 \times \text{ceiling} + 0.005 = 5.65 \times 10^{-3}$ at any evaluation point within the device-pass budget (the additive guard dominates at a floor-level ceiling, by design — PR-3 §B). The budget $B = 252,800$ device passes is twice the BPTT convergence point measured in a seed-7 sizing pilot excluded from the eight scored seeds.

A note this cell settles for free: the ceiling is *identical* under fixed-gain and saturating-gain substrate models ($\Delta_{M3} = 0$). The gain-model class — the largest modelling uncertainty in the substrate (§3) — cannot move the achievable accuracy, so it cannot flip any ranking below. The pre-registered M3 sensitivity trigger is therefore un-triggerable at this cell.

5.2 Gate ii: the recurrence trains on the device

Two decompositions of the Stage-0 gate were pre-registered (PR-9). **Capacity (ii-a):** the ceiling must clear a task-utility floor set at half the readout-only error — the reservoir baseline that freezes the recurrence and trains only the digital head. It clears it by $21.5 \times (5.2 \times 10^{-4})$ vs the reservoir’s 2.2×10^{-2} . **Trainability (ii-b):** at least one of the two hardware-committed routes — physics-aware training (PAT) or SPSA — must reach target on at least five of eight seeds.

Both reach target on all eight. This is the load-bearing result of the program: the parameters that *define* the recurrence — the per-ring detunings, the tunable ring–bus couplings, and the inter-ring couplings $\{\delta_j, \kappa_{\text{ext},j}, \mu_{jk}\}$ — are trained on the (simulated) physical substrate, through physical-operation-only gradient methods with fresh injected noise on every pass, to within the pre-registered margin of the exact-gradient ceiling. To our knowledge this is the first demonstration, by any method, of a recurrent photonic system whose recurrence-internal parameters are updated by on-device gradient-based or gradient-estimating training (the white-space claim, §1; [32]).

5.3 The ranking: sample-efficiency at matched device-pass cost

The primary metric is sample-efficiency: device passes to reach target, one physical pass being one sequence through the substrate in any direction (PR-7). Ranked lexicographically by success fraction then median passes (PR-8), with the digital-compute side-ledger co-reported but never folded into the rank:

route	success	median device passes → target	final SER (median)	digital ledger
PAT (twin-backward)	8/8	38,400	8×10^{-4}	505,600
adjoint (physical reverse pass)	8/8	73,600	5×10^{-4}	0
SPSA (model-free)	8/8	176,000	1.3×10^{-3}	0
RHEL (Hamiltonian echo)	0/8	censored at B	1.4×10^{-1}	0

All three ordered pairs among the passing routes separate with paired-by-seed bootstrap 95% confidence intervals excluding zero (PAT-SPSA = $-137,600$ passes, CI $[-155,200, -123,200]$; adjoint-PAT = $+35,200$, CI $[+35,200, +36,800]$; adjoint-SPSA = $-102,400$, CI $[-120,000, -88,000]$). The three routes trade the same axes the theory predicts they should. **PAT** is cheapest on the physical device but spends a **13×**-larger digital ledger and carries the full burden of characterizing a differentiable twin (§4, F8). **The adjoint** matches the exact-gradient ceiling in accuracy at zero digital cost and **1.9×** PAT’s device passes — though its count charges one physical reverse pass *as if* realizable, which no recurrent photonic system has yet demonstrated (a caveat we quarantine, §4). **SPSA** costs **4.6×** PAT’s device passes but needs no model, no twin, and no added hardware — the simplicity anchor of the Stage-1 plan.

No route earns promotion toward a later hardware slot (PR-9): the adjoint beats SPSA but loses to PAT on device passes, and clearing the bar requires clearly beating *both* workhorses. The hardware roadmap therefore stays on PAT and SPSA — the outcome the guardrail was built to protect, now settled by data rather than assertion.

5.4 RHEL under an honest echo

The fourth route, recurrent Hamiltonian echo learning (RHEL), is the one whose physical primitive SiN is least suited to supply. Rather than an idealized conjugation operator, we model the echo as a concrete $\chi^{(3)}$ four-wave-mixing phase-conjugation stage with its measured penalty chain — extraction, single-pass spiral conversion at 0.3 W pump, timing decay — totalling -22.4 dB per conjugation, plus the phase-insensitive parametric noise floor (§4, PR-11; [31]).

RHEL does not reach target on any seed; its final SER of **0.14** is *worse than the readout-only baseline* (a **+0.118** readout differential). Two controls locate the cause. A floor check confirms the estimator is correct: as the substrate is made progressively less dissipative, the RHEL gradient converges to the exact reference (direction cosine $\rightarrow 1.0000$; Fig. S5) — the non-dissipative limit RHEL’s theorem assumes. And an idealized-conjugator control at the smaller C-1 cell — a *perfect* echo, no conjugation loss or noise — *does* reach target ($5.7 \times 10^{-3} \leq 7.3 \times 10^{-3}$). So the failure at the headline cell is neither broken mechanics nor the conjugation chain: it is dissipative-echo bias, the irreducible mismatch between an echo that assumes time-reversal and a substrate that forgets its state within roughly nine of the sequence’s steps. This is the honest instantiation of the platform argument (§1, §5.4 of the proposal): silicon nitride’s low loss improves RHEL’s *noise* budget, but the *dissipation* the recurrence itself requires is fatal to the echo at the operating point. RHEL-on-SiN stays a simulation result.

5.5 The comparison the fair design was built to expose

One baseline result is more consequential for the program than the ranking. The offline-train-then-deploy route — train the full parameter set digitally on a designer’s model, then deploy through actuation maps, recalibrating only the digital head on-device — was given the *same* 5%-class calibration errors as PAT’s twin (the mismatch families are drawn from one frozen set, so the comparison cannot be rigged by giving in-situ training a secretly-wronger competitor; PR-5 F7.3). At that mismatch level it reaches 1.0×10^{-3} — statistically indistinguishable from in-situ PAT’s 8×10^{-4} .

We state the consequence plainly, because the fair-comparison design exists precisely to force it: **at 5% calibration accuracy on this task, training in situ buys essentially nothing over calibrate-then-deploy**. The demonstration claim — the first dissipative-resonator recurrence whose poles and couplings train on-device — stands regardless; it is a claim about *what was done*, not about beating an alternative. The question this raises — *under what conditions does the advantage appear?* — we then answered with two pre-registered follow-up experiments rather than leaving it open (PR-5 §E, PR-16; both frozen before the runs).

Calibration accuracy is not the axis (Fig. F8a). Sweeping the shared mismatch level from 5% to 30%-class (in-situ and offline drawing from one frozen family at every level, §5.1), in-situ training holds at the ceiling (7.8×10^{-4} , flat) while offline degrades only mildly ($1.0\text{--}1.2 \times 10^{-3}$) and *never fails the accuracy target* — the tie is robust to 30% calibration error. The differences are statistically real (paired-bootstrap CI excludes zero up to 20%) but never approach a factor of two: on this task the offline arm’s on-device head recalibration absorbs static parametric error, so a wrong recurrence with a well-fit head still equalizes. In-situ *recurrence* training does not earn its keep against calibration error alone.

Drift is the axis — specifically the part a re-lock cannot catch (Fig. F8b,c). We then let the substrate *drift*: a random walk on the ring detunings calibrated to a measured free-running silicon-nitride resonance drift (≈ 341 MHz over 24 h $\approx 24 \kappa_i$ at C-2 [33]), deployed after convergence, with each arm allowed its on-device response — offline recalibrates the head and re-locks the laser (a single global detuning re-centering); in-situ retraining the recurrence. The pre-registered contrast holds cleanly. Under **common-mode** drift (whole-chip thermal wander) the laser re-lock absorbs it and offline keeps pace (1.0 vs 0.8×10^{-3} , no advantage). Under **independent** per-ring drift the re-lock *cannot* fix the scrambled pole scatter, and in-situ retraining pulls ahead: 1.0×10^{-3} versus the re-locking offline’s 1.8×10^{-3} , a $1.84\times$ separation with paired-bootstrap CI excluding zero — and the gap *grows with accumulated drift* (reaching $\sim 3\text{--}4\times$ at the largest drift step). This $1.84\times$ time-integrated figure sits just below our pre-registered $2\times$ threshold, so we **do not declare a formal advantage** on the frozen metric; what the data show is a *statistically significant, mechanism-identified, drift-growing* in-situ edge that is specific to the uncorrelated drift a global lock leaves behind. A larger drift magnitude or an end-of-deployment metric would cross the threshold, but neither was pre-registered, so we name them rather than claim them.

Reporting a null where the design was built to detect an advantage is the honest core of §5.5, not a hedge around it: in-situ training’s value here is not calibration robustness (null to 30%) but tracking the uncorrelated drift a laser lock cannot — and quantifying *how uncorrelated real on-chip drift is* becomes the sharpest Stage-1 measurement (§9).

5.6 What the diagnostics add

Two mechanism rows, at three seeds each on C-1, support the mismatch narrative without inflating it (Fig. S2). Decomposing PAT’s twin mismatch — perfect twin, parametric-error twin, structural-omission twin (dropping the gain self-consistency channel) — all three reach the C-1 ceiling identically: PAT absorbs both mismatch families at this cell. We flag explicitly that this does **not** extrapolate to C-2, where the dropped gain channel was measured to carry $\sim 8\%$ of the gradient direction (§4, the adjoint cosine dropping from 0.994 to 0.925 with cell size); the C-2 mismatch sensitivity is a Stage-0.5-full measurement, not an inference from C-1. And the idealized-RHEL row is the control cited in §5.4. Neither row is a headline; both are the pre-registered controls that let the headlines mean what they say.

5.7 Controllability: what “N = 32” actually means

The bake-off’s cell label understates a constraint that any hardware implementation inherits, so we report it as a first-class result (Fig. F3). Under a single input tap, the per-ring gradient magnitude collapses geometrically with distance from the drive — by ring 32 it sits some twenty-five orders of magnitude below the maximum — and the *participation profile* (settled per-ring amplitude relative to the maximum) counts only $\{1, 3, 5\}$ of 32 rings above $\{10^{-1}, 10^{-2}, 10^{-3}\}$. A nominally 32-ring lattice driven at one port is, effectively, a three-ring computer with 29 passengers. This is the in-data form of the program’s reservoir-falsifier: it is *why* the readout-only baseline stalls at 2.2×10^{-2} (§5.2), and it is a controllability property of the chain physics, not of any training method.

The pre-registered remedy is a measured, minimal input map: the smallest tap set (capped at $K = 4$) under which *every* ring’s gradient clears 10^{-3} of the maximum. The resolved map, taps $\{3, 12, 21, 30\}$, clears the gate for all 32 rings with worst ratio 1.42×10^{-3} — taken as a *minimum over five drive realizations*, because single-seed margins at the gate boundary flicker by a factor of ~ 20 . No three-tap set clears (best: 24 of 32), and the registered starting guess $\{1, 9, 17, 25\}$ was not the winner (28 of 32 — its worst ring sat seven hops from a tap). Under the resolved map the participation counts rise to $\{4, 26, 32\}$, and every use of “ $N = 32$ ” in this paper carries that measured profile rather than the nominal dimension. Two protocol rulings made during resolution are on record: the gate references the *maximum-gradient* ring (the registered ring-1 reference is gameable when ring 1 is untapped), and robustness binds on the min-over-seeds. Both strengthen the gate; both were adopted before the finalists were evaluated. The price of controllability is charged honestly where it lands: four drive E/O channels instead of one, priced in the systems envelope (§7) — trainability of the deep lattice is bought with exactly the conversion overhead the advantage question (§7.4) must then carry.

5.8 Secondary diagnostic (registered, deferred)

PR-14 — the bias/variance decomposition of each estimator’s gradient against the BPTT reference — is registered as a secondary diagnostic only and was not run in the core bake-off; it is a Stage-0.5-full row. The reason it is secondary is structural: gradient-direction agreement flatters the exact methods (adjoint, RHEL) and penalizes SPSA, whose per-step alignment is poor by construction while its *averaged* trajectory converges (§5.3) — scoring on cosine alone would have reproduced the known failure mode of ranking estimators by a proxy the task does not pay for. The fragments that exist (the adjoint’s 0.994/0.925 cell-dependent cosine, §5.6; RHEL’s non-dissipative-limit recovery, §5.4) are reported where they carry mechanistic weight, and the full decomposition belongs in an appendix when the S0.5-full rows run.

6. Choosing the damping operating point

The D-LinOSS line’s central observation — that damping in an oscillatory SSM is a *performance knob*, not a defect to minimize — has a sharp physical meaning here: per-ring damping is the net loss $\kappa_{\text{net}}(\kappa_{\text{ext}})$, and the tunable coupler that sets it is already in the trained partition. The registered disposition (R-ii) therefore frames the question not as “which damping value do we freeze?” but as “does training *find* the right damping within the feasible box?” — and the experiment separates the two readings with two arms, six damping points, eight seeds each, at the headline cell under the convergence-controlled protocol of §5.1.

The damping value matters enormously. With κ_{ext} *pinned* (only detunings and inter-ring couplings training), the converged error spans a factor of ~ 300 across the feasible box: near-chance performance at light damping ($r = 0.2$: SER 0.39) falling to 1.3×10^{-3} at the optimum — which sits at **deep overcoupling** ($r^* = 2.0$, $\kappa_{\text{net}} \approx 4.1\kappa_i$, about 5.5 samples of memory at 2 GS/s). The direction is instructive: the equalization task needs only a 7-tap span, and the long memory the light-damping regime supplies (30–45 samples) is actively harmful — stale symbols interfere. “More memory” is not free performance in a dissipative recurrence; damping tunes memory *to the task*, which is precisely the D-LinOSS thesis in physical units, with the optimum on the heavily-damped side for this task class.

Training absorbs the knob. In the second arm the full partition trains inside progressively wider boxes $[r_{\text{min}}, r_{\text{hi}}]$. Every box containing the pinned optimum reaches within the pre-registered margin of it — in fact reaching 5×10^{-4} , the §5 ceiling, *better* than the best uniform pin: per-ring trainable coupling finds a heterogeneous damping profile no single design value can express. Boxes that exclude the good regime fail exactly as they must (training cannot find what the clamp forbids). The registered R-ii test is therefore **confirmed**:

the designer’s damping obligation is to make the feasible box *contain* the good regime; the operating point itself is the trained substrate’s job.

Two honest footnotes. The slow mid-grid configurations ($r = 0.3\text{--}0.5$) had not fully plateaued at the training ceiling, so the $\times 300$ spread is a budget-bounded statement — but both endpoints and all winning configurations converged, and the R-ii verdict uses converged points only. And the result recolors §5 slightly: a large share of what in-situ training accomplished in the bake-off *is* finding the damping operating point (the θ_0 hold value, pinned, yields 0.038 — seventy-seven times worse than the trained substrate). Since the offline baseline finds r^* on its calibrated model just as well (§5.5), the damping result strengthens the trainability story without moving the advantage question.

7. Does it pay? The systems envelope

7.1 The question, and how we keep it honest

A photonic SSM only matters if, after paying the conversion toll — DAC and modulator in, photodiode and ADC out, heaters held all the while — it still beats a competent digital implementation on the axis the niche cares about. We price this at two frozen corners (OPT = best published device class; CONS = named vendor parts at ENOB-at-speed, never nominal bits), against named baselines (Microsoft Brainwave’s author-stated batch-1 streaming efficiency; a coherent-DSP ASIC class; Jetson AGX Orin), with every number traced to a frozen source row and the four known exclusions (laser wall-plug, locking, control compute, packaging) explicitly unbudgeted — they only shrink positive cells, so negative findings are robust to them. An assembly-time retrieval bounds their magnitude from primary sources (supplementary ledger): the integrated-class stack — hybrid-laser wall-plug 0.2–0.6 W at mW-class on-chip power, TEC hold 0.18 W, microcontroller-class control 0.13 W, FPGA-class locking up to ~ 10 W — totals of order 0.5–1 W, the *same order as the budgeted* $N = 128$ photonic power itself (~ 0.9 W at 2 GS/s). Charging it would compress the positive cells’ margin (§7.2) toward single digits, and a benchtop realization (40–100 W laser and instrument lock) would erase the niche outright. The niche verdict below therefore carries an integrated-realization condition on all four excluded items, alongside the heater-class condition it already states.

7.2 Inference: a conditional niche, gated by the heater class

The lite envelope’s verdict stands: **a plausible low-latency niche exists, conditionally**. At the registered scale grid, the photonic side clears the strongest streaming baseline (Brainwave) by up to $\sim 15\times$ in energy per sample at $N = 128$ and 2 GS/s, with end-to-end latency bounded at tens of nanoseconds per sample against the baseline’s milliseconds — but only when three conditions hold simultaneously: line rates ≥ 0.5 GS/s (every scenario loses everything at 0.1 GS/s), $N \geq 32$ (conversion is N -independent; digital cost scales with N — the structural effect the architecture banks on), and **suspended low-power heaters** (~ 1 mW/ π , class B). Under the registered worst-case holding convention the foundry-standard heater class loses to every baseline everywhere in the window at the deployable corner; the demonstrated-foundry path therefore does not reach the energy niche as computed — a Stage-1 platform constraint stated as such, with the expected-value-holding sensitivity (under which the optimistic corner clears in-window) reported alongside per the review finding. Jetson’s *peak* rating is never beaten anywhere; the niche claim rests on measured sustained behavior, for which batch-1 recurrent workloads on edge GPUs are documented at $>100\times$ below claimed peak — page-verified: measured batch-1 GRU throughput of 1.9 and 3.5 GOP/s against claimed peaks of 0.5 and 0.8 TOP/s on the two Jetson-class devices (ratios $\approx 263\times$ and $\approx 229\times$), with the source’s own conclusion stating “a factor of over 100X” [34]; the Orin DLA path falls back to GPU for recurrent layers. One boundary cell (the DSP-class comparison at $N = 32$) clears by $\sim 1\%$ and is treated as a tie.

7.3 Training: the energy metric inverts the sample-efficiency ranking

The bake-off’s primary metric (§5) counts device passes, under a registered principle that the digital side-ledger is always co-reported and never merged. The envelope is where that principle pays off, because converting both ledgers to joules **inverts the ranking**. Training to the pre-registered target at the headline cell costs, at the optimistic corner:

route	conversion energy	digital compute	total
PAT	0.57 mJ	$\approx 13\text{--}44\text{ J}$ (twin ledger, 1.3×10^{13} FLOP at named accelerator classes)	$\approx 13\text{--}44\text{ J}$
adjoint	1.1 mJ	0	1.1 mJ (†realizability)
SPSA	2.6 mJ	0	2.6 mJ

PAT reaches target in the fewest device passes but its digital twin bill exceeds SPSA’s *entire training energy* by roughly four orders of magnitude — and our FLOP estimate is charitable to PAT (real accelerator utilization on a 64-dimensional complex recurrence sits far below peak). SPSA — no model, no twin, no added hardware — trains the physical recurrence to target for **about 2.6 millijoules, all-in**, at the optimistic corner (99 mJ at the vendor-part corner). RHEL’s echo, censored on accuracy grounds anyway, is also energy-dominated by its own conjugator pump ($\times 42$ its conversion stack per update). For Stage 1 this sharpens §9’s ordering: SPSA is not merely the simplest route but by far the cheapest to *run as training*, and PAT’s role is best cast as the high-device-throughput option for settings where digital compute is free and device time is scarce — which is a real regime (a shared testbed), but a different claim than “efficient.”

7.4 The advantage question, answered as far as the data allows

Assembling §5.5, §6, and this section: the *demonstration* is in-data; the *advantage* is conditional and partly open. In-situ training buys nothing over calibrate-then-deploy at 5%-class calibration accuracy (§5.5) — and much of what training achieves on this task is finding the damping operating point, which a well-calibrated offline model also finds (§6). The inference-mode niche exists but is gated by a heater class the named foundry flow does not supply, and rests on sustained-vs-peak baseline conventions we document rather than hide. What survives all of it: a recurrent photonic SSM at GS/s line rates with class-B actuation is energy-competitive at scale for streaming workloads, can be *trained through its own physics for millijoules* when calibration is unavailable or stale, and offers latency headroom no digital baseline in our set approaches. Whether the conjunction of those conditions describes a market or only an experiment is a Stage-1 question, and §9 designs the experiment to answer it.

8. Limits of this model

8.1 Robustness here means robustness to what we modelled

Every robustness statement in §5 is conditioned on the substrate of §3: the methods absorb the imperfections *we simulated* — saturating gain at a registered operating point, Langevin ASE at NF 7 dB, resolved backscatter doublets, 5%-class calibration mismatch, fresh noise per pass. Hardware contains channels we did not model: thermal transients and self-heating at operating power, polarization rotation, fabrication disorder beyond the derate row, drift at cadences between our episode and training scales, and mode-splitting behavior that is not captured by a single always-on γ per process class. Any of these could reorder the §5 ranking on a real chip. We regard the ranking as a hypothesis the Stage-1 hardware exists to test, with PAT and SPSA committed precisely because they are the two routes whose chip-level robustness is already literature-established [3, 2].

8.2 Anchor risks and verification debts, by name

The substrate’s realism leans on anchors with stated residual risks: the C-2 loss class transfers a wide-multimode racetrack result to a single-mode registry ring (priced by the $\times 2$ -loss derate row, §3.3); the Er:Si₃N₄ noise budget rests on a **single coupling-loss-limited measured NF (~7 dB)** in the flagship device paper — debt #3’s original “no measured NF” premise was found false at the S0.3-0 recon and our NF-A = 7.0 dB was frozen to match the measurement (§3.2); the narrower residue (intrinsic amplifier NF not isolated) keeps the NF sensitivity rows registered; the recurrent adjoint pass is charged *as if realizable* with no demonstration in the literature (debt #4 — an inference from absence, time-stamped mid-2026); and the white-space claim itself is one-sided evidence from a pre-registered search, to be re-swept before submission (debt #1). The S0.4-0 calibration retired one internal debt (the drive build-up controversy resolved by measurement: $\times 0.42$, doublet-quenched) and left one open: the on-resonance floor calibration under worst-case de-saturation (anchor-risk vii, §3.6), carried as a label.

8.3 The benchmark anchor we do not use

An early gate required reproducing a published LinOSS benchmark as an external anchor. Running the authors’ own code on their published seeds reproduced their headline within noise on one dataset but not on the long-sequence EigenWorms task, where we traced a numerical-precision failure mode in the training loss (an absorbing zero-gradient state in fp32) that makes the published number seed-unstable (Fig. S1; a separate reproducibility note is in preparation). The gate was adjudicated purpose-served-with-anchor-void: all downstream accuracy references in this program are therefore **in-house BPTT-on-substrate ceilings** measured under our own protocol (§5.1), never transferred published numbers. We flag fp32-sensitivity generally: our substrate runs float64, and the eval-floor granularity of §5 is symbol-count-limited, not precision-limited.

8.4 Review independence

Through 2026-07-06 every freeze in the ledger passed adversarial review by an independent reviewer session reporting to the PI, and several results in this paper exist because that review forced them (the multi-tap input map, the mismatch decomposition, the readout-differential rule). From 2026-07-07 the program ran in a single-session mode in which the same agent performed both execution and review, under standing PI delegation; every artifact from that period is so labelled in the ledger, and the S0.4b/c/S0.5 findings — including the two honest nulls (the offline tie; RHEL’s failure) — should be read with that reduced independence in mind. The pre-registration discipline (thresholds frozen and committed before runs) is the structural mitigation; it is auditable in the supplementary commit trail regardless of who held the pen.

8.5 Scope limits we chose

The task family is deliberately narrow (continuous-signal channel equalization plus a synthetic memory family; the registered secondary task is deferred), and the C-3 128-ring cell never gates anything. Two axes the bake-off itself held fixed were measured afterward in pre-registered follow-ups (§5.5): calibration-mismatch sensitivity (swept to 30%-class — the tie holds) and drift (a literature-calibrated random walk under a deploy-then-drift protocol, two correlation regimes). Drift remains unmodelled *during* training at the bake-off cadence, and the tested drift magnitude is gentle ($\approx 1.4 \kappa_i$ accumulated) rather than worst-case. The systems-advantage question — whether any of this pays once conversion overhead is counted — is §7’s; the strongest current evidence is §5.5’s sharpened verdict: no advantage from calibration accuracy at any tested mismatch, and a statistically significant but sub-threshold edge specific to uncorrelated drift. We consider stating that plainly to be the paper’s job.

9. Outlook: Stage 1

9.1 What the first chip must be

The bake-off fixes the Stage-1 chip's training stack by evidence rather than taste: **PAT and SPSA, nothing else in the loop**. Neither exotic route earned promotion (§5.3), and the hardware ledger (supplementary F8) shows why that is unlikely to reverse on hardware grounds alone: the adjoint adds circulators, phase-coherent reverse injection, and an unsolved separation of the counter-propagating field from the backscatter doublet; RHEL adds a pumped conjugator bank whose power budget (≈ 9.6 W at 32 rings) exceeds the entire rest of the system. SPSA's row is the quiet asset — zero added components, zero model burden — so the minimal viable demonstration is: the §3 plant ($N = 8\text{--}32$ rings, foundry-floor Q suffices per C-1's gate), thermo-optic $\{\delta, \kappa_{\text{ext}}, \mu\}$ actuation, one drop-port readout chain, the four-tap drive map of §3.6, and SPSA as the first-light training route with PAT layered on once the twin is characterized to the 5%-class the mismatch protocol assumed.

The multi-project-wafer path is concrete: the registered cells were chosen to be foundry-realizable (C-1 at generic-foundry loss; C-2 bounded by a demonstrated MPW result [20]), and the actuation map of §2 uses only standard thermo-optic tuners. The E/O overhead the four-tap drive and the eval protocol add is exactly what §7's envelope prices.

9.2 What would change our mind

Three pre-registered forks, with their triggers on record: (i) **adjoint promotion** — if a Stage-1-adjacent demonstration retires debt #4 (a physical recurrent reverse pass), the PR-9 criterion re-opens with the S0.5 data as prior; the sim says it would arrive at ceiling-grade accuracy at $2\times$ PAT's device cost, zero digital. (ii) **The in-situ advantage** — the offline tie (§5.5) at 5%-class mismatch sets the burden: in-situ training earns its place on hardware only where calibration is worse than 5%-class, drifts on deployment timescales, or where the envelope penalizes the twin's digital ledger. The Stage-1 experiment should be *designed to measure exactly this differential* — same chip, offline-deploy vs PAT/SPSA arms — rather than assume it. (iii) **RHEL** — nothing on SiN; the sim verdict (dissipation-fatal at the operating point even with a perfect conjugator) would need a *conservative* platform regime, not a better conjugator, to reopen.

9.3 Beyond the linear unselective core

The frozen architecture is deliberately the unselective LTI core — poles and couplings, the part photonics builds natively. The selectivity axis (input-dependent dynamics in the Mamba direction [35]) maps onto the same lattice as input-dependent $\mathbf{C}$ then $\mathbf{B}$ actuation and is scoped for a later stage only behind its own gate (per-step tuning without per-state DACs); nothing in this paper's claims depends on it. Likewise the damping operating point (§6) and the D-LinOSS accuracy question ride the *trainable* κ_{ext} axis established here rather than new hardware.

9.4 Closing

The program set out to answer a narrow question with unusual bookkeeping: can the physics of a dissipative photonic recurrence be trained through itself, and at what honest cost? In simulation, under pre-registered thresholds: yes — by the two methods a chip can already run, at device-pass costs now quantified, with the exotic routes priced out by data and the advantage-over-offline question left honestly open. The next sentence worth writing requires a chip.

Figures

(a) trained partition $\{\delta_j, \kappa_{\text{ext},j}, \mu_{jk}\}$, per-ring saturating gain g (b) realizable memory region




Figure F1 — Architecture and realizable pole region. (a) The $N = 32$ coupled-ring chain (C-2 cell): trained parameter set $\{\delta_j, \kappa_{\text{ext},j}, \mu_{j,j+1}\}$ (ring detunings, bus couplings, inter-ring couplings), with the resolved four-tap input map $B = \{3, 12, 21, 30\}$ (§3, §5.7). (b) Realizable memory (samples at 2 GS/s) versus intrinsic Q at gain fractions $g_f \in \{0, 0.5, 0.9\}$, the three registered cells (C-1 foundry-floor, C-2 headline, C-3 aspirational), and the 7-tap task-span line — why the operating point carries gain: at $g_f = 0$ the foundry-floor cell sits at the task-span line; at the registered $g_f = 0.9$ all three cells clear it with margin.




Figure F2 — The dissipative substrate's operating map. (a) Settled saturating net loss $\kappa_{\text{net}}(r)$ versus the fixed-gain plane, the lasing crossing r^* , the registered clamp band ($r_{\text{min}} = 0.1606$ with margin $m_{\kappa} = 0.05$, $\Delta r = 0.02$), and the initialization point θ_0 . (b) Off-resonance de-saturation at r_{min} versus θ_0 — the quantified anchor-risk (vii): a detuned ring at r_{min} reaches $\kappa_{\text{net}} = -0.48 \kappa_i$, which is why the clamp is referenced on-resonance (§3.6).




Figure F3 — Controllability is a first-class constraint. Per-ring gradient magnitude (relative to the maximum ring) under a single input tap versus the resolved four-tap map. The single-drive profile collapses geometrically with distance from the drive (ring 32 sits ~ 25 orders below the maximum; participation $\{1, 3, 5\}/32$ rings above $\{10^{-1}, 10^{-2}, 10^{-3}\}$); the resolved $B = \{3, 12, 21, 30\}$ puts all 32 rings above the pre-registered 10^{-3} gate (worst ring 1.42×10^{-3} , min over five drive seeds).




Figure F4 — Sample efficiency to target (the headline). Median held-out SER versus physical device passes at C-2 (8 seeds; shaded IQR): PAT, recurrent adjoint, SPSA, RHEL (honest echo, censored at budget), the readout-only reservoir baseline, and the offline-deploy arm (deployed pre-trained, so it starts low; head recalibration only).

Dotted lines mark the BPTT-on-substrate ceiling (device passes = 0 by the PR-7 convention, drawn as a level only) and the pre-registered target SER = 5.65×10^{-3} ; the vertical line is the device-pass budget $B = 252,800$.




Figure F5 — Ranking at matched device-pass cost. Median device passes to target (per-seed dots) with the digital-computation side-ledger co-reported (hatched; PAT’s twin backward = 505,600 digital passes): PAT 38,400 < adjoint 73,600 < SPSA 176,000; RHEL censored 0/8. The ranking answers the pre-registered promotion question in the negative: neither exact method beats both workhorses.




Figure F6 — Damping is a first-order design knob. Final SER versus uniform pinned overcoupling r (plateaued endpoints spanning $\times 302$), the deep-overcoupling optimum $r^* = 2.0$ ($\kappa_{\text{net}} \approx 4.1 \kappa_i$ — excess memory is harmful

for this task), and the trainable- κ_{ext} box (R-ii): boxes containing r^* train to the 5×10^{-4} ceiling, beating every uniform pin — the heterogeneous damping profile is found by training, not designed.




Figure F7 — The systems envelope. (a) Training-energy inversion: SPSA trains the C-2 cell all-in for ~ 2.6 mJ (optimistic conversion accounting; 99 mJ conservative) while PAT's device side is 0.57 mJ but its digital twin backward costs 13–44 J — the energy metric inverts the device-pass ranking. (b) Inference energy per sample versus state dimension N at 2 GS/s against the named digital baselines: the photonic envelope (optimistic corner, low-power heater class) clears the measured embedded-GPU (Jetson sustained) and Brainwave batch-1 lines at all N and enters the DSP-ASIC class at $N = 128$ — while never beating the Jetson *peak-spec* line, which we report alongside: the niche is conditional, as §7 states.




Figure F8 — What breaks the offline tie (pre-registered follow-ups, §5.5). (a) Calibration-mismatch sweep, 5→30%-class (8 seeds): in-situ PAT holds at the ceiling (7.8×10^{-4} , flat) and offline-deploy degrades only mildly, never failing target — crossover $m^* = \text{none}$; the tie is robust to calibration error. (b) Deploy-then-drift, common-mode regime ($\sigma_{\text{step}} = 0.40 \kappa_i$ per step on all detunings coherently): the offline laser re-lock absorbs the drift and keeps pace with in-situ retraining. (c) Independent per-ring drift: the re-lock cannot fix per-ring pole scatter; in-situ retraining holds near-ceiling while the re-locking offline baseline degrades with accumulated drift (time-integrated ratio 1.84 $\times$, paired-bootstrap CI excluding zero — below the pre-registered 2 $\times$ advantage threshold, so reported as a directional edge, not a declared advantage). Dashed/dotted lines: target and ceiling.




Figure S1 — The benchmark anchor we do not use (G3 dossier, §8.3). (a) Validation-accuracy trajectories of the *official* LinOSS-IM code on the published EigenWorms seeds (our rerun, 2026 stack): two of five seeds visibly collapse mid-training (the fp32 absorbing-zero-gradient mode). (b) Final test accuracy per seed against the published $95.0 \pm 4.4\%$: rerun mean 90.56% , $\sigma 9.34 \approx 2.1 \times$ the published dispersion (per-seed $97.22 / 83.33 / 97.22 / 97.22 / 77.78$).




Figure S2 — Twin-mismatch decomposition at C-1 (§5.6). Final SER (3 seeds) for PAT with a perfect twin, a parametric-error (M-par) twin, and a structural-omission (M-struct) twin — all three at the C-1 BPTT ceiling, i.e. mismatch channels ≈ 0 at *this cell only* (the dropped gain channel grows to $\sim 8\%$ of gradient direction at C-2, §5.6) — and the idealized-conjugator RHEL control, which reaches the C-1 target and thereby localizes RHEL's C-2 failure to echo physics, not mechanics.




Figure S3 — The concrete echo sub-model (PR-11, §4). (a) Conjugation-chain waterfall at the frozen operating point (mechanism A, shared spiral bank): ring-port extraction η_{ex}^2 , single-pass χ^3 -FWM spiral conversion (0.3 W pump, 0.5 m), routing/insertion — chain $\eta_c = -22.4$ dB per conjugation. (b) Per-cell feasibility ceilings for the alternative mechanisms: resonant-ring loaded-Q ceiling from the state bandwidth (mechanism B) and 10-ns off-chip transit amplitude survival (mechanism C).




Figure S5 — RHEL recovers its own theorem's limit (§5.4). Cosine between the RHEL update and the exact BPTT gradient as the substrate is made progressively less dissipative: -0.75 at $\kappa_{net}Tdt \approx 1.0$, monotonically to $+1.0000$ at 0.03 — exact recovery of the non-dissipative limit. The registered C-2 operating point sits at $\kappa_{net}Tdt \approx 27$, far beyond the anti-aligned regime: the C-2 failure is dissipative-echo bias, not implementation error.

References

1. Spall, “Multivariate stochastic approximation using a simultaneous perturbation gradient approximation,” *IEEE Trans. Autom. Control* 37(3), 332–341 (1992)
2. Bandyopadhyay, Sludds, Krastanov, Hamerly, Harris, Bunandar, Streshinsky, Hochberg, Englund, “Single-chip photonic deep neural network with forward-only training,” *Nat. Photonics* 18, 1335–1343 (2024). arXiv:2208.01623
3. Wright, Onodera, Stein, Wang, Schachter, Hu, McMahon, “Deep physical neural networks trained with backpropagation,” *Nature* 601, 549–555 (2022)
4. Hughes, Minkov, Shi, Fan, “Training of photonic neural networks through in situ backpropagation and gradient measurement,” *Optica* 5(7), 864–871 (2018)
5. Tanaka et al., “Recent advances in physical reservoir computing: A review,” *Neural Networks* 115, 100–123 (2019); Van der Sande, Brunner, Soriano, “Advances in photonic reservoir computing,” *Nanophotonics* 6(3), 561–576 (2017)
6. Jayatilaka et al., “Wavelength tuning and stabilization of microring-based filters using silicon in-resonator photoconductive heaters,” *Opt. Express* 23(19), 25084–25097 (2015); Mak, Sacher, Xue, Mikkelsen, Yong, Poon, “Automatic Resonance Alignment of High-Order Microring Filters,” *IEEE JQE* 51(11) (2015); Milanizadeh, Aguiar, Melloni, Morichetti, “Canceling Thermal Cross-Talk Effects in Photonic Integrated Circuits,” *JLT* 37(4), 1325–1332 (2019)
7. Bueno, Maktoobi, Froehly, Fischer, Jacquot, Larger, Brunner, “Reinforcement learning in a large-scale photonic recurrent neural network,” *Optica* 5(6), 756–760 (2018)
8. Böhm, Verschaffelt, Van der Sande, “A poor man’s coherent Ising machine based on opto-electronic feedback systems...,” *Nat. Commun.* 10, 3538 (2019); companion: Böhm et al., *Nat. Commun.* 13, 5847 (2022)
9. Wu et al., “Monolithically integrated asynchronous optical recurrent accelerator,” *eLight* 5, 7 (2025)
10. Ashtiani, Idjadi, Kim, “Integrated photonic neural network with on-chip backpropagation training,” *Nature* 651, 927–932 (2026). arXiv:2506.14575
11. Zhao et al., “In-Situ Trained Microring-Based Neural Networks,” *Laser Photon. Rev.* (2025), 10.1002/lpor.202501576
12. Wu, Ren et al., “Time-synthetic optical neural networks with stable programmable gain,” arXiv:2507.02297 (retitled from “A scalable and programmable optical neural network in a time-synthetic dimension”)
13. Zhang, Wang, Lederman, Shastri, Prucnal et al., “Compact, reconfigurable, and scalable photonic neurons by modulation-and-weighting microring resonators,” *eLight* 6, 6 (2026). arXiv:2505.11369
14. “In-situ optimization of an optoelectronic reservoir computer with digital delayed feedback,” *ACS Photonics* (2025). arXiv:2502.11126
15. “Real-time optical signal equalization with a silicon photonic spatially distributed reservoir computer,” *Nat. Photonics* (2026). arXiv:2503.19911 (Ghent/imec)
16. Gu, Goel, Ré, “Efficiently Modeling Long Sequences with Structured State Spaces,” *ICLR 2022*. arXiv:2111.00396
17. Gu, Gupta, Goel, Ré, “On the Parameterization and Initialization of Diagonal State Space Models,” *NeurIPS 2022*. arXiv:2206.11893
18. Rusch, Rus, “Oscillatory State-Space Models,” *ICLR 2025 (Oral)*. arXiv:2410.03943

19. Boyer, Rusch, Rus, “Learning to Dissipate Energy in Oscillatory State-Space Models,” arXiv:2505.12171
20. Cui, Cao, Pan, Gao, Yu, Zhang, “Compact microring resonator based on ultralow-loss multimode silicon nitride waveguide,” *Adv. Photonics Nexus* 2(4), 046007 (2023)
21. Gupta, Gu, Berant, “Diagonal State Spaces are as Effective as Structured State Spaces,” NeurIPS 2022. arXiv:2203.14343
22. Haus, *Waves and Fields in Optoelectronics*, Prentice-Hall (1984)
23. Gorodetsky, Pryamikov, Ilchenko, “Rayleigh scattering in high-Q microspheres,” *JOSA B* 17(6), 1051–1057 (2000); Kippenberg, Spillane, Vahala, “Modal coupling in traveling-wave resonators,” *Opt. Lett.* 27(19), 1669–1671 (2002)
24. Liu, Huang, Wang, He, Raja, Liu, Engelsen, Kippenberg, “High-yield, wafer-scale fabrication of ultralow-loss, dispersion-engineered silicon nitride photonic circuits,” *Nat. Commun.* 12, 2236 (2021)
25. Rukh, Colación, Buck, Drake, “Process-structure-property relationships in subtractive fabrication of silicon nitride microresonators for nonlinear photonics,” arXiv:2511.02198 (2025)
26. Liu, Qiu, Ji, Lukashchuk, He, Riemensberger, Hafermann, Wang, Liu, Ronning, Kippenberg, “A photonic integrated circuit-based erbium-doped amplifier,” *Science* 376, 1309–1313 (2022)
27. Ye, Marpaung, “Compact multi-mode silicon-nitride micro-ring resonator with low loss,” *Adv. Photonics* 5(5), 050503 (2023)
28. Talandier, multi-layer photonic-equalization manuscript (in preparation) + chip β -track thermal-crosstalk/perturbation-tolerance results
29. Pourcel, Ernout, “Learning long range dependencies through time reversal symmetry breaking,” arXiv:2506.05259 (2025)
30. López-Pastor, Marquardt, “Self-Learning Machines Based on Hamiltonian Echo Backpropagation,” *Phys. Rev. X* 13, 031020 (2023). arXiv:2103.04992
31. Riemensberger, Kuznetsov, Liu, He, Wang, Kippenberg, “A photonic integrated continuous-travelling-wave parametric amplifier,” *Nature* 612, 56–61 (2022); Krückel et al., “Continuous wave-pumped wavelength conversion in low-loss silicon nitride waveguides,” *Opt. Lett.* 40(6), 875–878 (2015)
32. this work: PR-15 two-modality search + refresh memos (docs/s0_L/debt1_whitespace_search.md, docs/s0_L/whitespace_refresh_2026-07-12.md), supplementary
33. Dacha, Zhao, McNulty, Bhatt, Lipson, Gaeta, “Frequency-stable nanophotonic microcavities via integrated thermometry,” *Nature Photonics* (2025). arXiv:2506.21692
34. Gao, Rios-Navarro, Chen, Liu, Delbruck, “EdgeDRNN: Recurrent Neural Network Accelerator for Edge Inference,” *IEEE JETCAS* 10(4), 419–432 (2020). arXiv:2012.13600
35. Gu, Dao, “Mamba: Linear-Time Sequence Modeling with Selective State Spaces,” COLM 2024 (Outstanding Paper). arXiv:2312.00752